



You've got the lighter, let's start a fire: Spark and Cassandra

Robert Stupp

Solutions Architect @ DataStax, Committer to Apache Cassandra

@snazy

Cassandra Days brought to you by DataStax



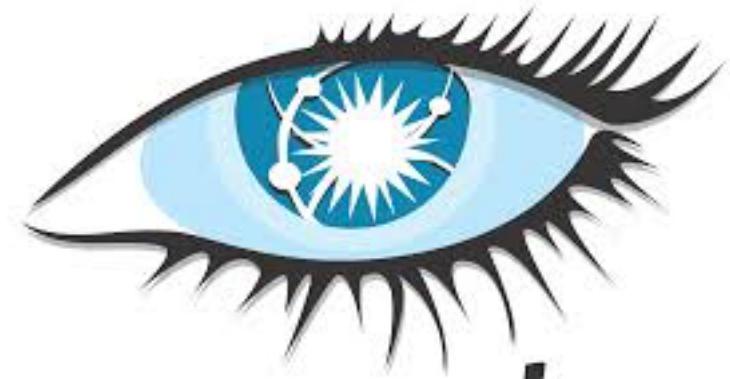
A journey through
Distributed Computing



KillrPower Application

Show current power consumption in Germany in real time

- Power consumption of all cities on a map
- Graph of data for one city



cassandra



One word...



THE
APACHE[®]
SOFTWARE FOUNDATION

Spark™

The Kafka logo features a black icon of interconnected circles of varying sizes, resembling a network or a cluster. To the right of the icon, the word "kafka" is written in a lowercase, sans-serif font.

Apache Cassandra

Distributed & resilient database



Cassandra Use-Cases

- Time Series
- Personalization
- (or a mix of both)



Cassandra?

- Joining data ... hmm
- Analyzing the whole data set ... hmm
- Ad-Hoc Queries (SQL) ... hmm



Apache Spark

Distributed & resilient analytics

*“Spark lets you run programs up to 100x faster in memory,
or 10x faster on disk, than Hadoop.”*



Spark : Simplify the...

challenging and compute-intensive task of processing high volumes of real-time or archived data, both structured and unstructured,

seamlessly integrating relevant complex capabilities such as machine learning and graph algorithms.



Source: <http://www.toptal.com/spark/introduction-to-apache-spark>

Spark Use Cases

- Event detection
- Behavior / anomalies detection
- Fraud & intrusion detection
- Targeted advertising
- E-commerce transaction processing
- Recommendations



Spark Use Cases

- Data migration
- Ad hoc queries
- Business Intelligence



Spark

Spark
SQL

Spark
Streaming

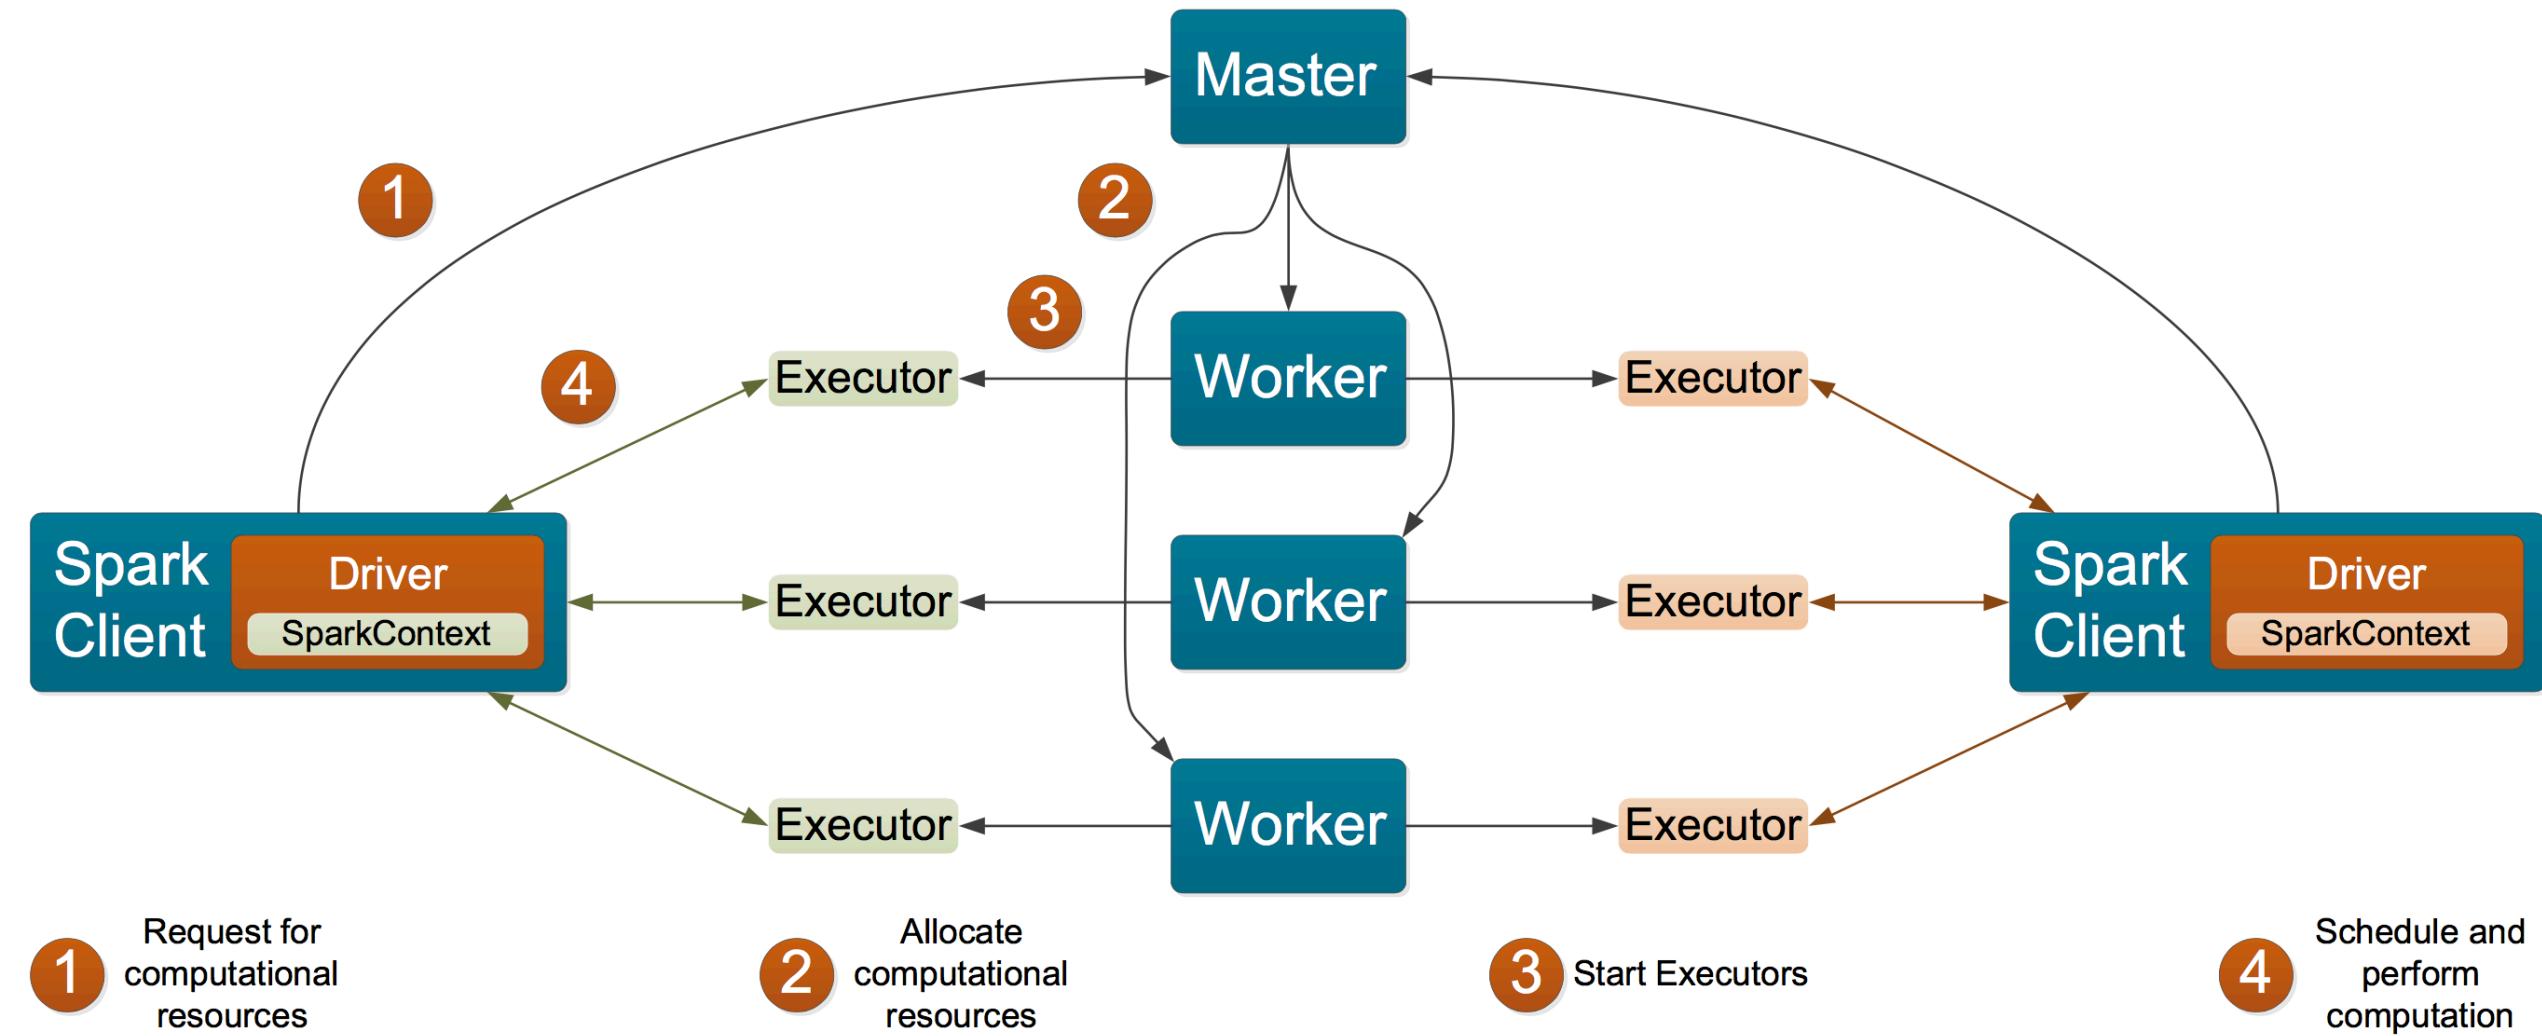
MLlib

GraphX

SparkR



Spark Architecture



Spark RDDs

Resilient Distributed Datasets

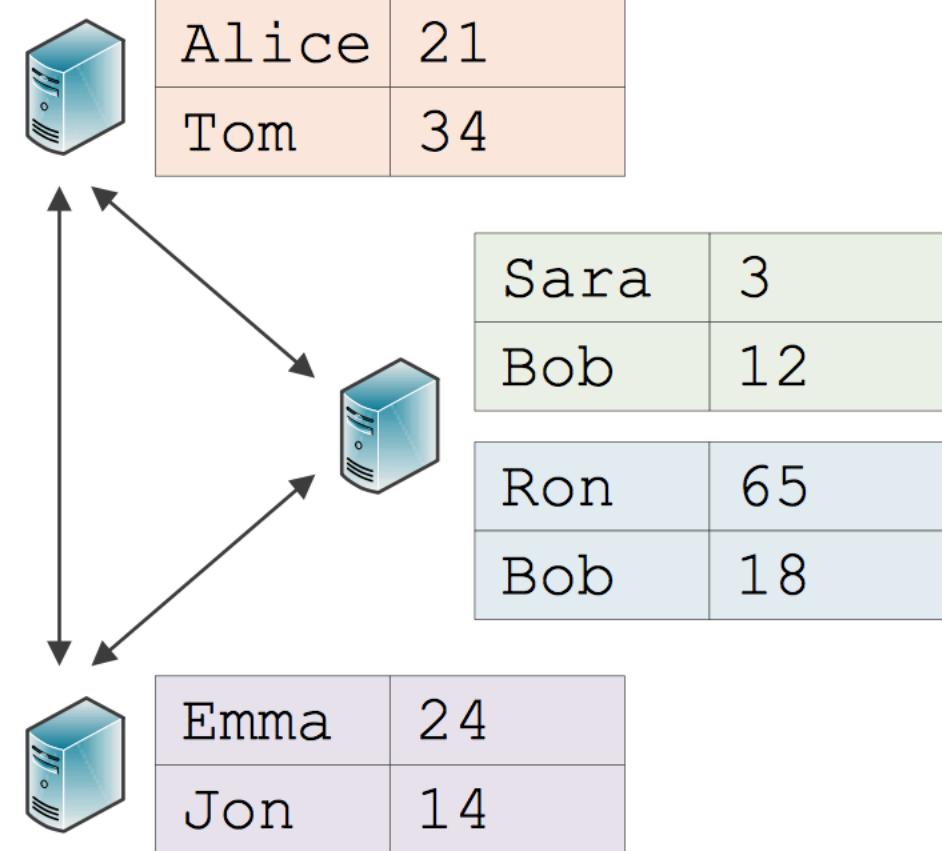
- Fault-tolerant collection of elements
- Can be operated on in parallel
- Immutable
- In-Memory (as much as possible)



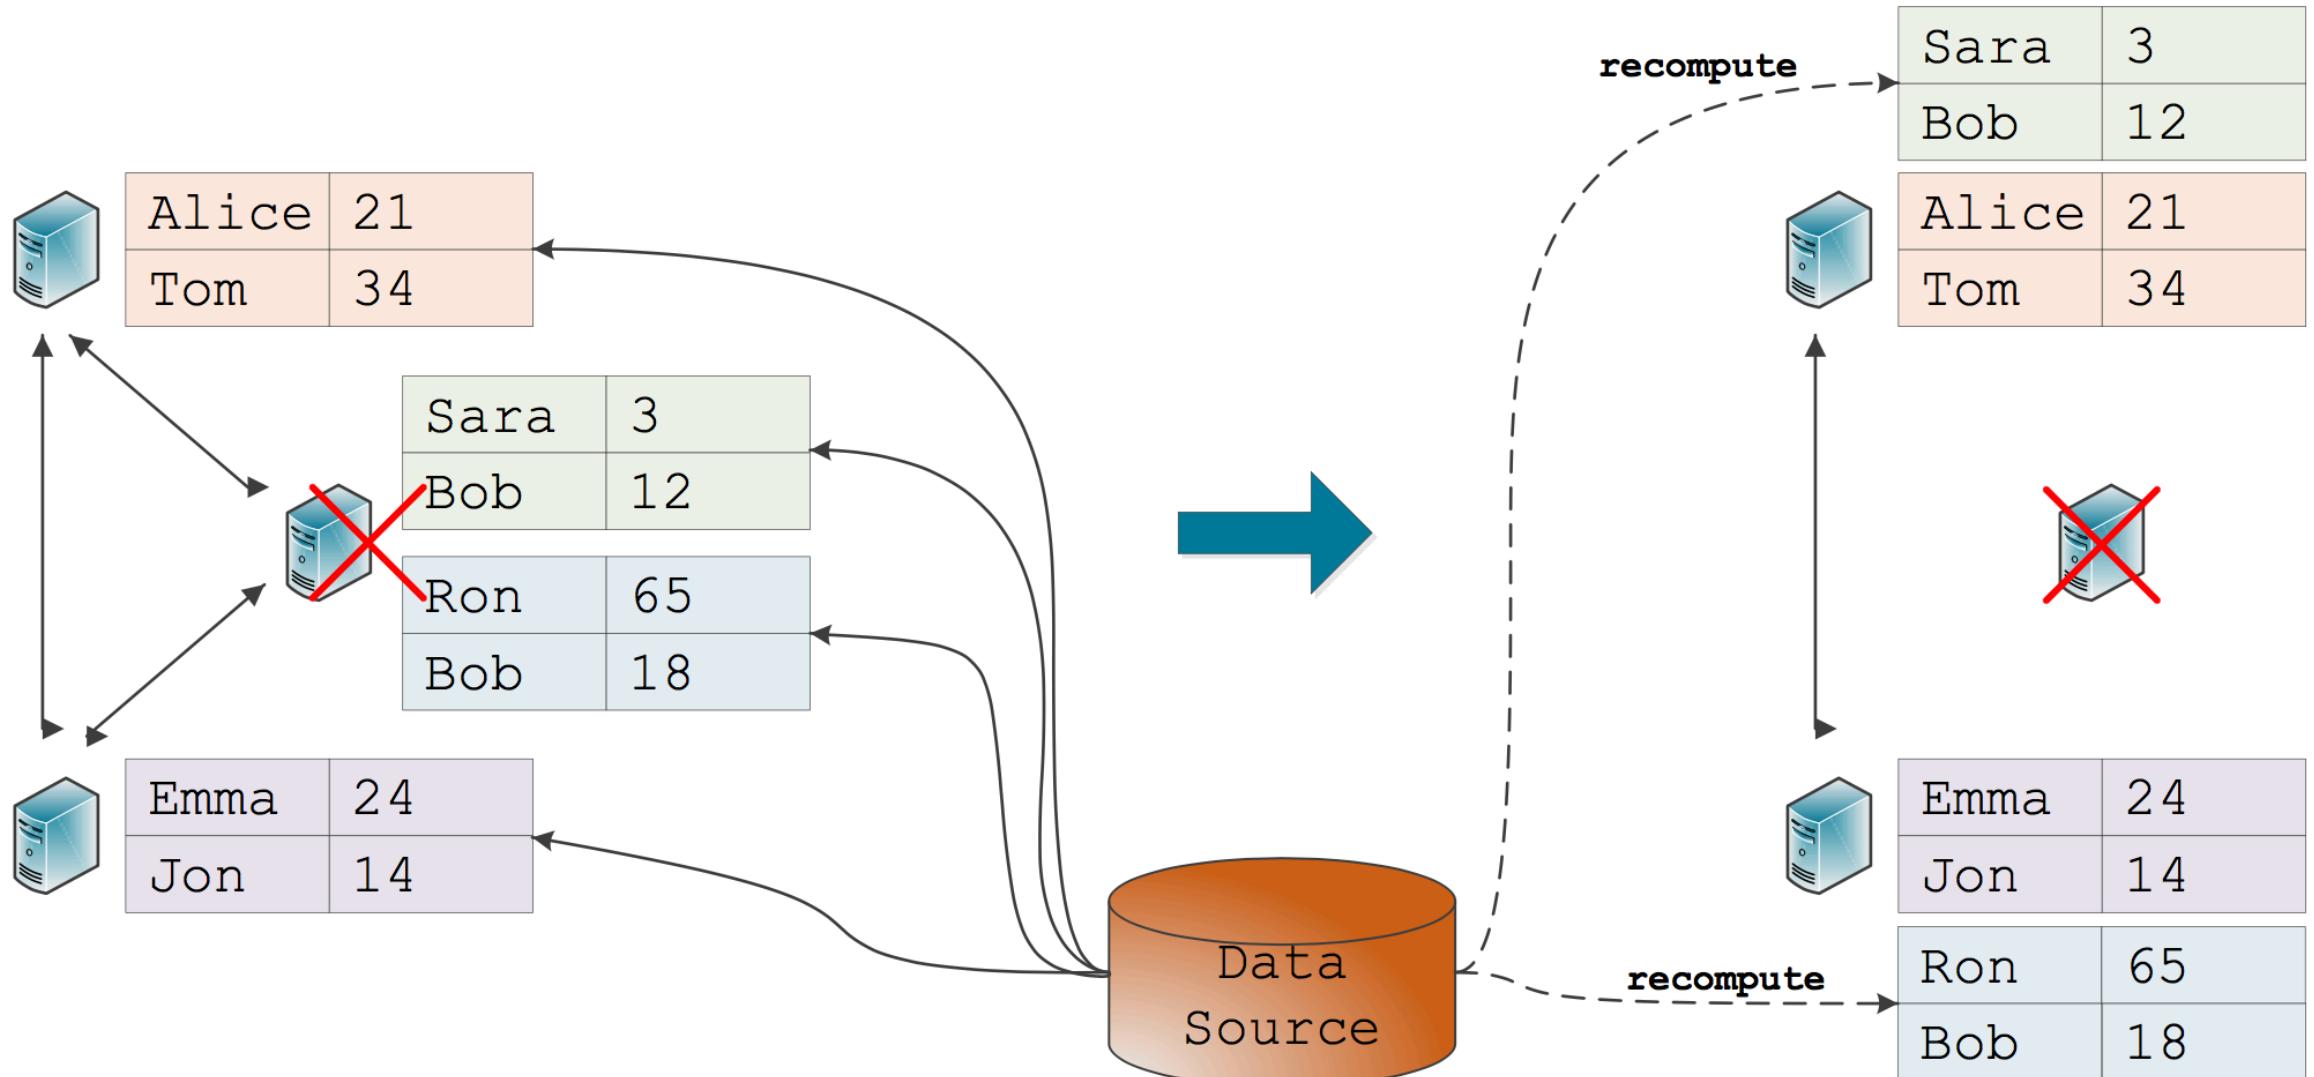
Spark RDD Partitioning

RDD[(String, Int)]

String	Int
Alice	21
Tom	34
Sara	3
Bob	12
Ron	65
Bob	18
Emma	24
Jon	14



Spark RDD Resiliency



Spark Transformations

Create a new data set from an existing one

Transformation	Meaning
<code>map(func)</code>	Return a new distributed dataset formed by passing each element of the source through a function <i>func</i> .
<code>filter(func)</code>	Return a new dataset formed by selecting those elements of the source on which <i>func</i> returns true.
<code>flatMap(func)</code>	Similar to map, but each input item can be mapped to 0 or more output items (so <i>func</i> should return a Seq rather than a single item).
<code>mapPartitions(func)</code>	Similar to map, but runs separately on each partition (block) of the RDD, so <i>func</i> must be of type <code>Iterator<T> => Iterator<U></code> when running on an RDD of type T.
<code>mapPartitionsWithIndex(func)</code>	Similar to mapPartitions, but also provides <i>func</i> with an integer value representing the index of the partition, so <i>func</i> must be of type <code>(Int, Iterator<T>) => Iterator<U></code> when running on an RDD of type T.

Spark Actions

Return a value after running a computation

Action	Meaning
<code>reduce(func)</code>	Aggregate the elements of the dataset using a function <i>func</i> (which takes two arguments and returns one). The function should be commutative and associative so that it can be computed correctly in parallel.
<code>collect()</code>	Return all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.
<code>count()</code>	Return the number of elements in the dataset.
<code>first()</code>	Return the first element of the dataset (similar to <code>take(1)</code>).
<code>take(n)</code>	Return an array with the first <i>n</i> elements of the dataset. Note that this is currently not executed in parallel. Instead, the driver program computes all the elements.

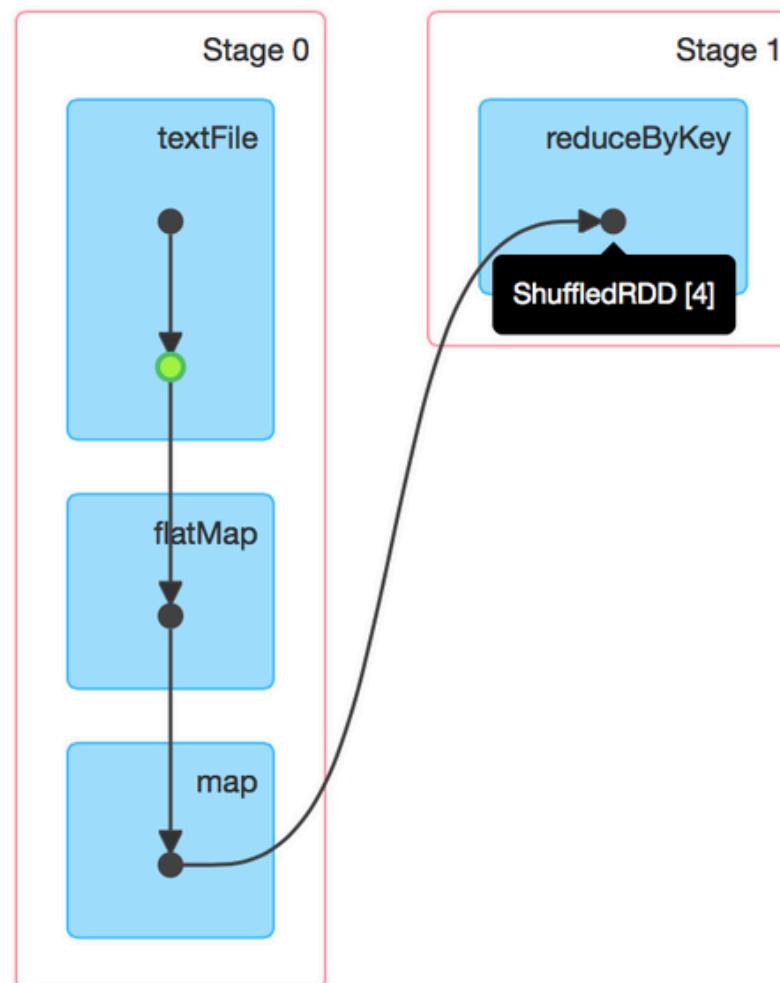
Spark DAG

Details for Job 0

Status: SUCCEEDED

Completed Stages: 2

- ▶ Event Timeline
- ▼ DAG Visualization



Spark Example (word count)

Given is a text file with movies and their categories.

We want to know the occurrences of each genre.

Pirates of the Caribbean, Adventure, Action, Fantasy

Star Trek I, SciFi, Adventure, Action

Die Hard, Action, Thriller



Why Scala

Scala is a functional language

Analytics is a functional “problem”

So – use a functional language



Spark Example (word count)

```
sc.textFile("file:///home/videos.csv")
    .flatMap(line => line.split(",").drop(1))
    .map(genre => (genre,1))
    .reduceByKey{case (x,y) => x + y}
    .collect()
    .foreach(println)
```



Spark Example (word count)

```
sc.textFile("file:///home/videos.csv")
    .flatMap(line => line.split(",").drop(1))
    .map(genre => (genre, 1))
    .reduceByKey{ case (x, y) => x + y}
```

Anonymous Scala Function



Spark Example (word count)

```
sc.textFile("file:///home/videos.csv")
    .flatMap(line => line.split(",").drop(1))
    .map(genre => (genre,1))
    .reduceByKey{case (x,y) => x + y}
    .collect()
    .foreach(println)
```

Flat the collection



Spark Example (word count)

```
sc.textFile("file:///home/videos.csv")
    .flatMap(line => line.split(",").drop(1))
    .map(genre => (genre,1))
    .reduceByKey{case (x,y) => x + y}
    collect()
```

Pair RDD



Spark Example (word count)

```
sc.Values of TWO Pair RDDs
    .textFile("hdfs://node1:9000/videos.csv")
    .flatMap(line => line.split(",").drop(1))
    .map(genre => (genre, 1))
    .reduceByKey{case (x,y) => x + y}
    .collect()
    .foreach(println)
```



Spark Example (word count)

```
sc.textFile("file:///home/videos.csv")
    .flatMap(line => line.split(",").drop(1))
    .map(genre => (genre,1))
    .reduceByKey{case (x,y) => x + y}
    .collect()
    .foreach(println)
```



Spark Example (word count)

```
sc.textFile("file:///home/videos.csv")
    .flatMap(line => line.split(",").drop(1))
    .map(genre => (genre,1))
    .reduceByKey{case (x,y) => x + y}
    .collect()
    .foreach(println)
```



Cassandra-Spark-Driver

Spark-Cassandra-Connector

- Scala, Java
- Open Source (ASF2)
- Developed by DataStax and community guys

[https://github.com/datastax/
spark-cassandra-connector](https://github.com/datastax/spark-cassandra-connector)



Spark Batch Processing

Other REALLY useful stuff like

- Re-partition data & processing
- Allows efficient use of Cassandra



SQL and Data Frames

Spark SQL

```
csc.sql(" SELECT COUNT(*) AS total          " +
         " FROM killr_video.movies_by_actor " +
         " WHERE actor = 'Johnny Depp'      ")
.show
```

```
+----+
|total|
+----+
|   54|
+----+
```



Spark Data Frames

```
val movies = csc
  .read
  .format("org.apache.spark.sql.cassandra")
  .options(Map( "keyspace" -> "killr_video",
    "table" -> "movies_by_actor" ))
  .load
```

```
movies.filter("actor = 'Johnny Depp' ")
  .agg(Map( "*" -> "count" ))
  .withColumnRenamed( "COUNT(1)", "total" )
  .show
```

Same as:

```
SELECT COUNT(*) AS total
FROM killr_video.movies_by_actor
WHERE actor = 'Johnny Depp'
```



Spark Data Frames

Data frame is RDD that has named columns

- Has schema
- Has rows
- Has columns
- Has data types



Spark Streaming

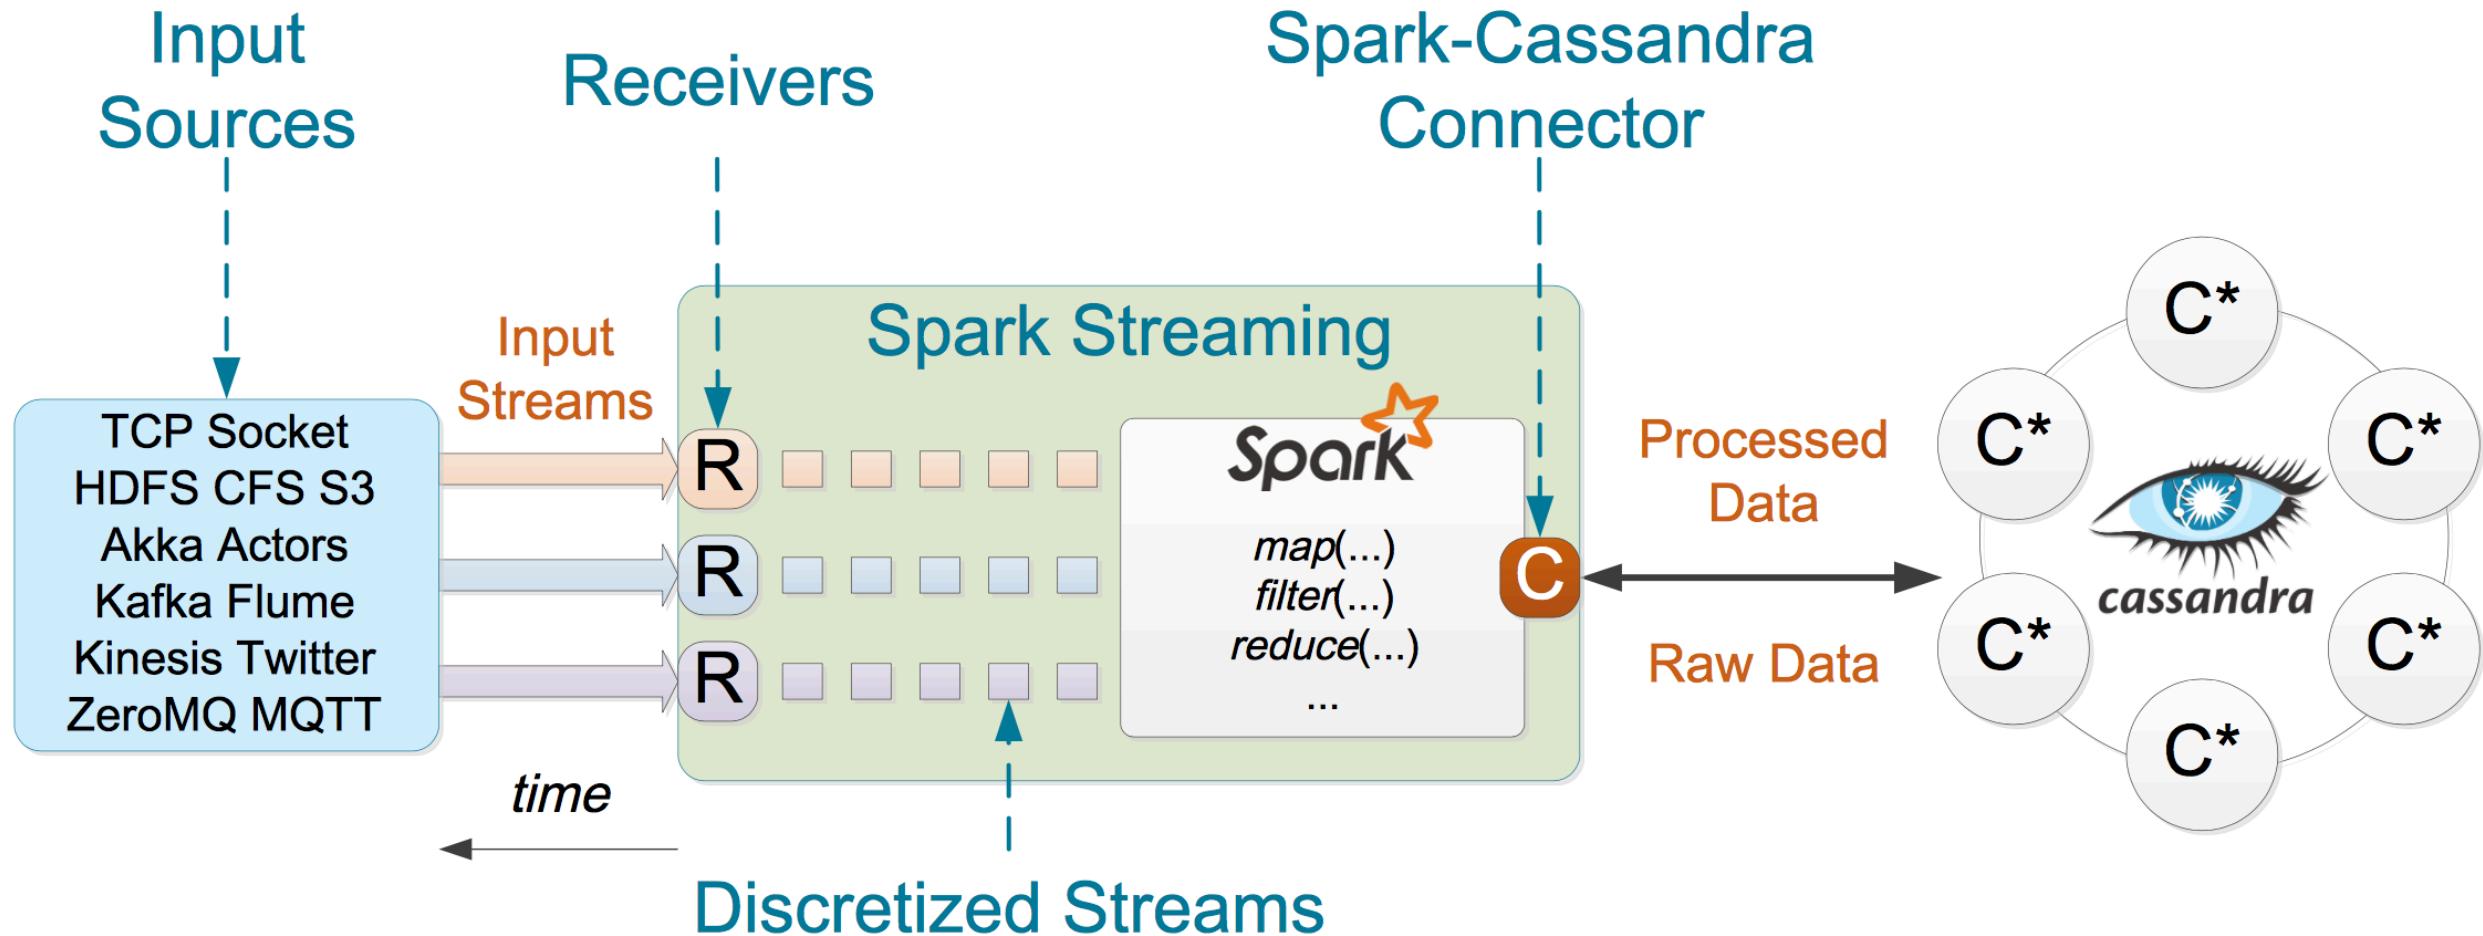
Spark Streaming

Continuous Micro Batching

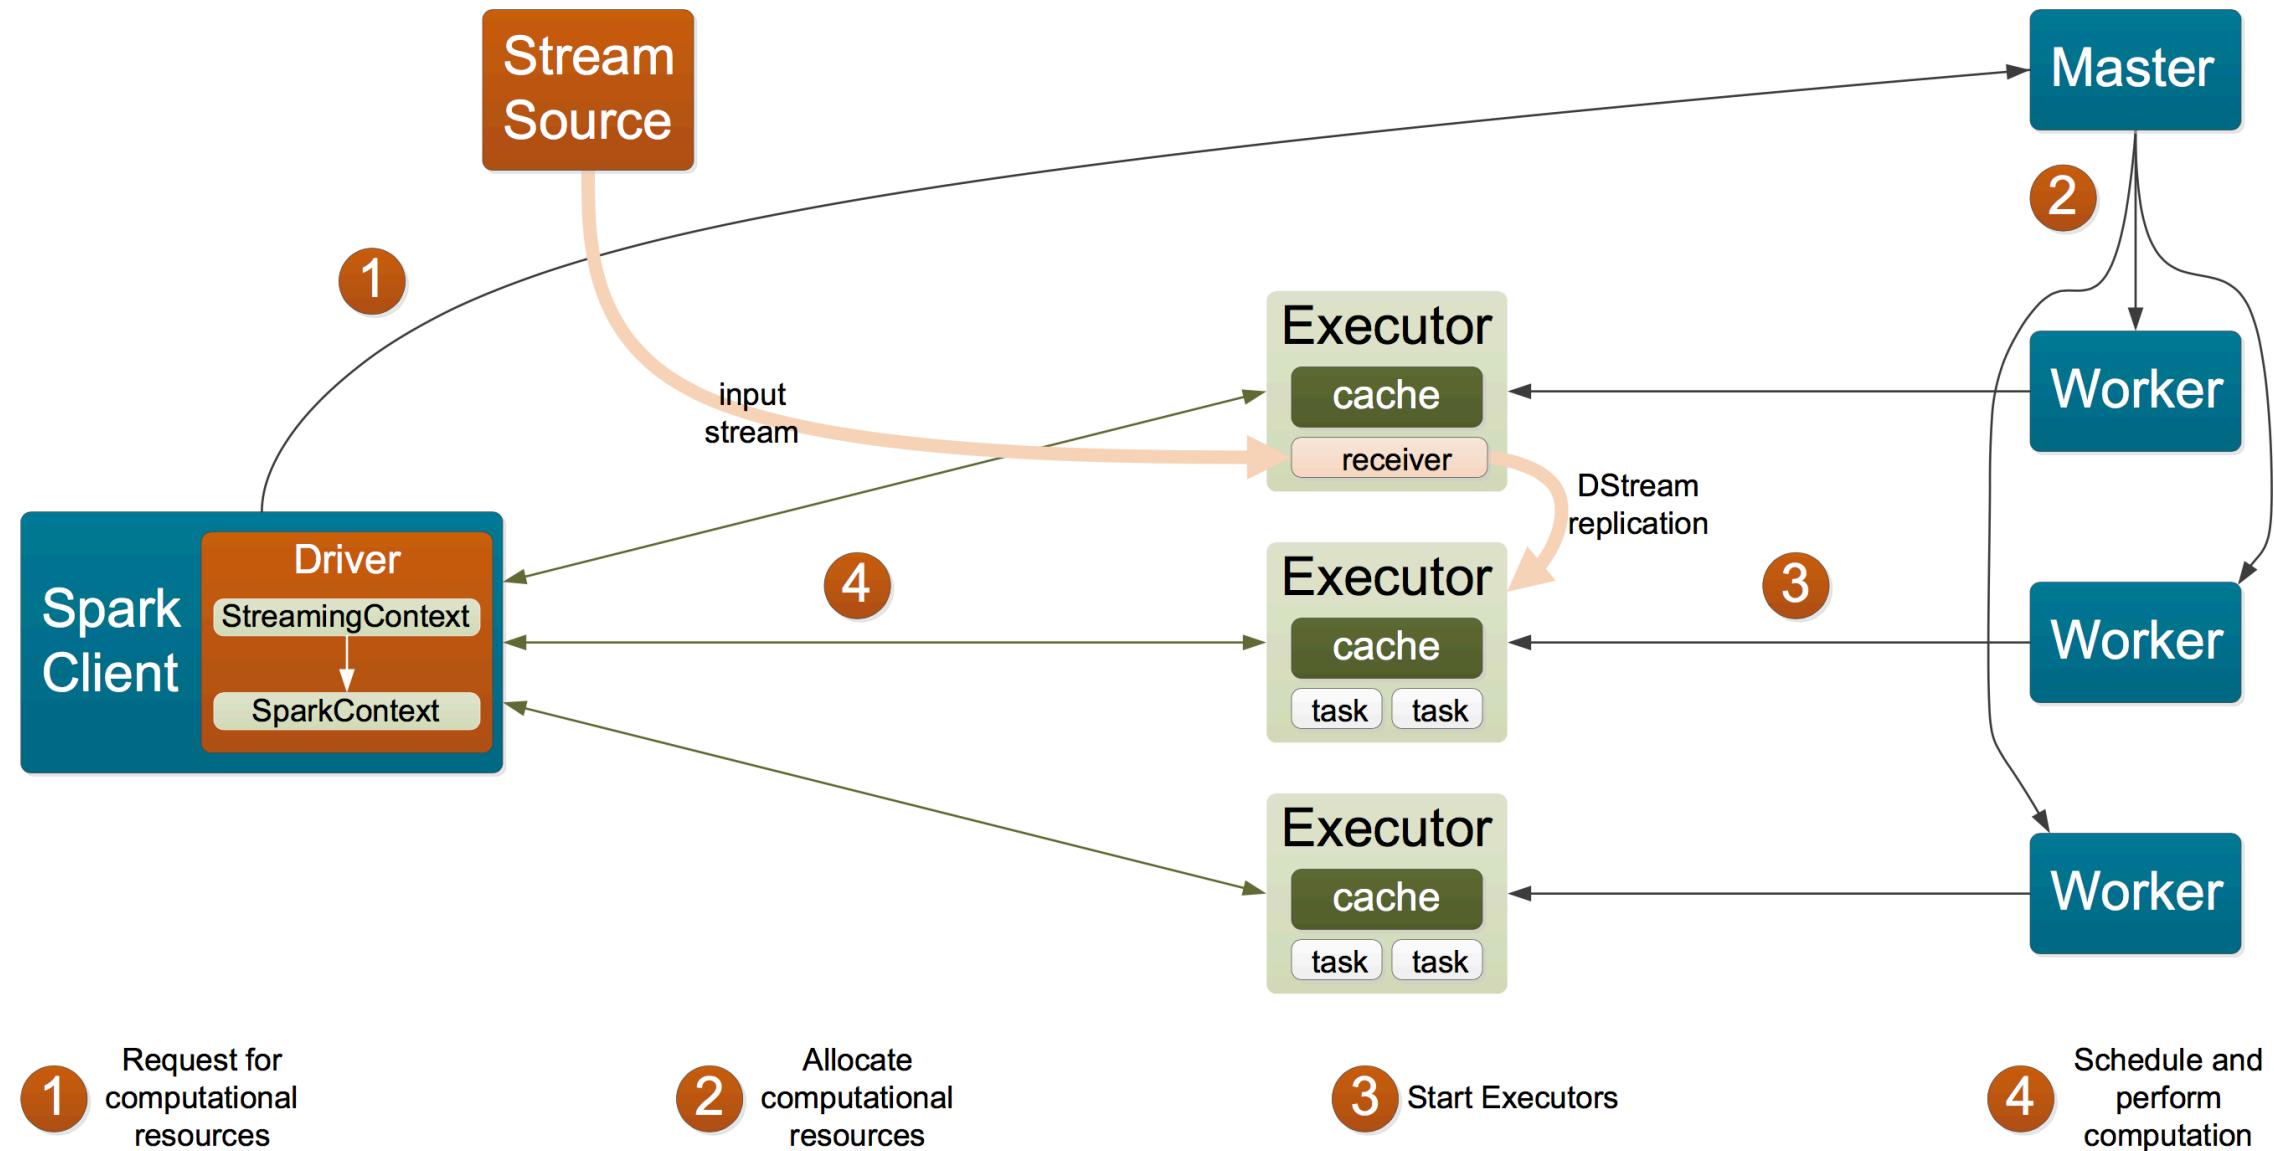
- Collect data of time slides
- Compute on that data



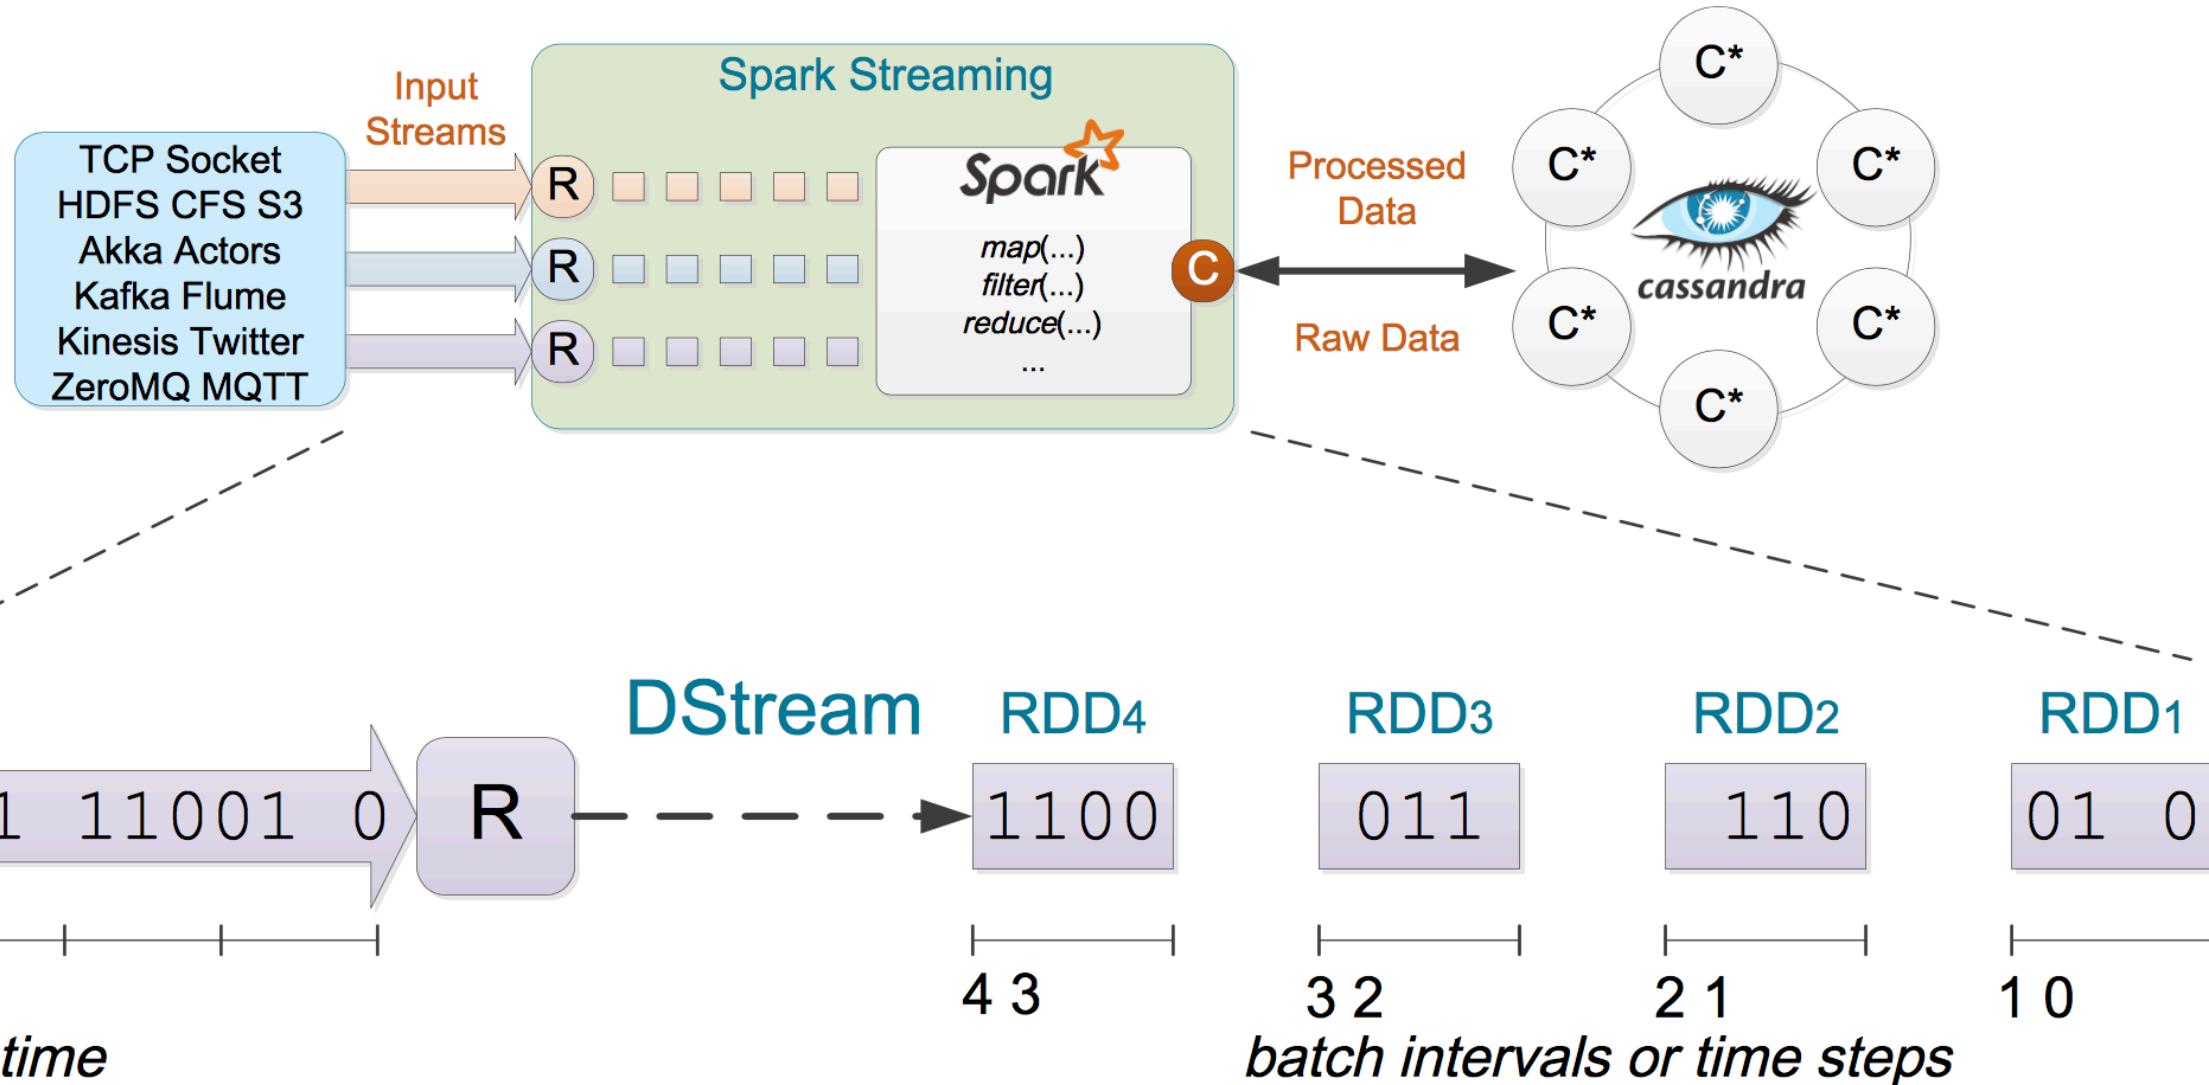
Spark Streaming



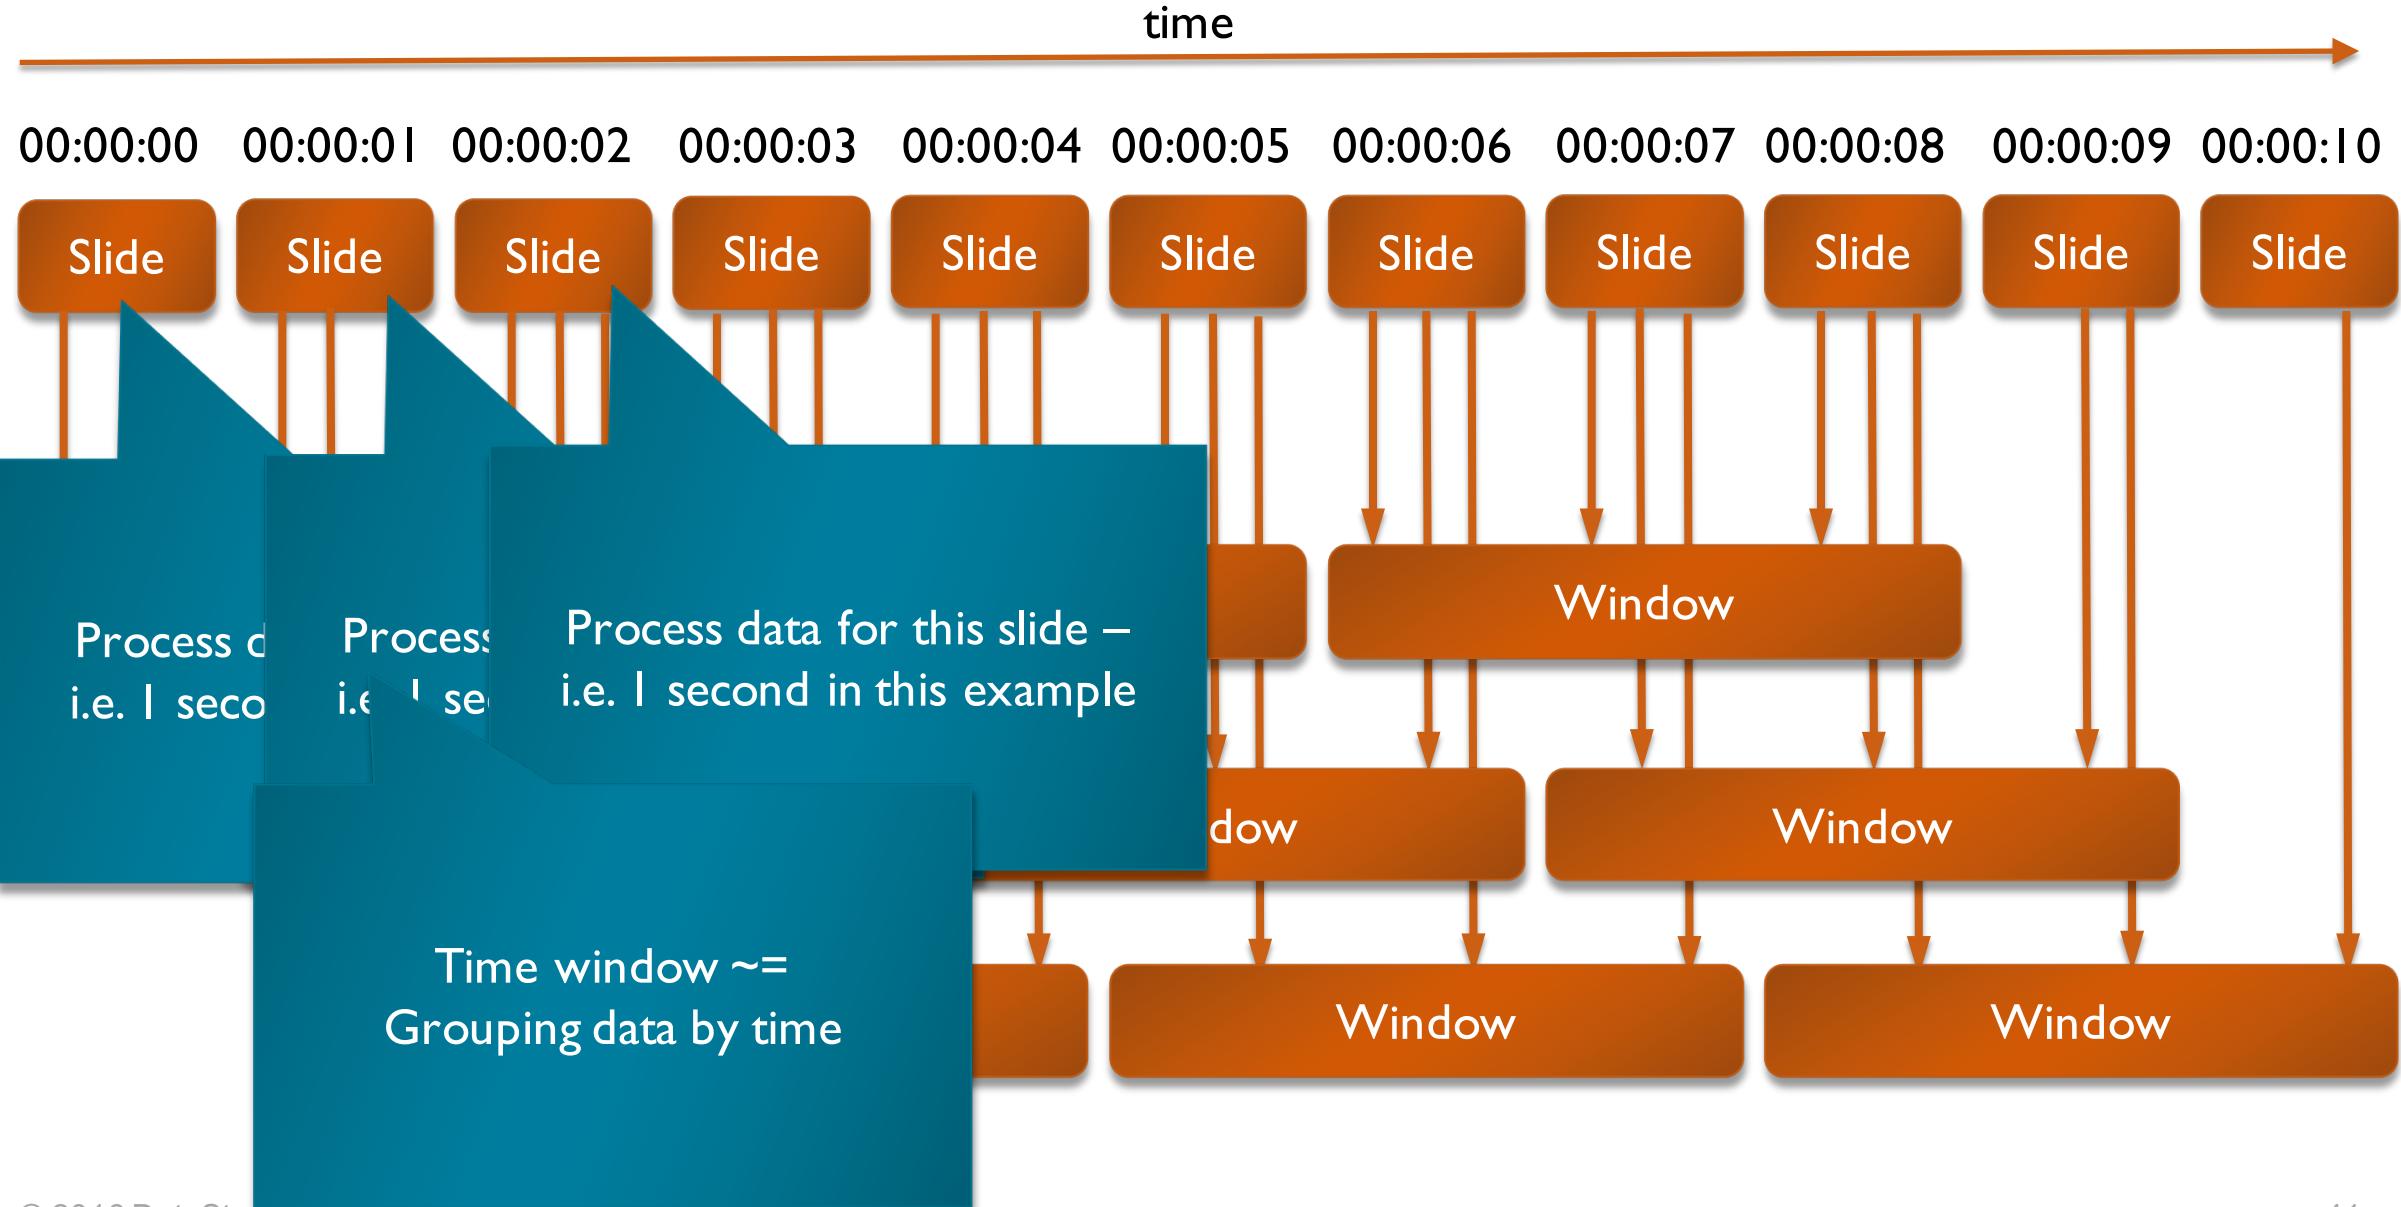
Spark Streaming Architecture



Spark Streaming



Spark Streaming Windows



Apache Kafka

Distributed & **resilient** messaging

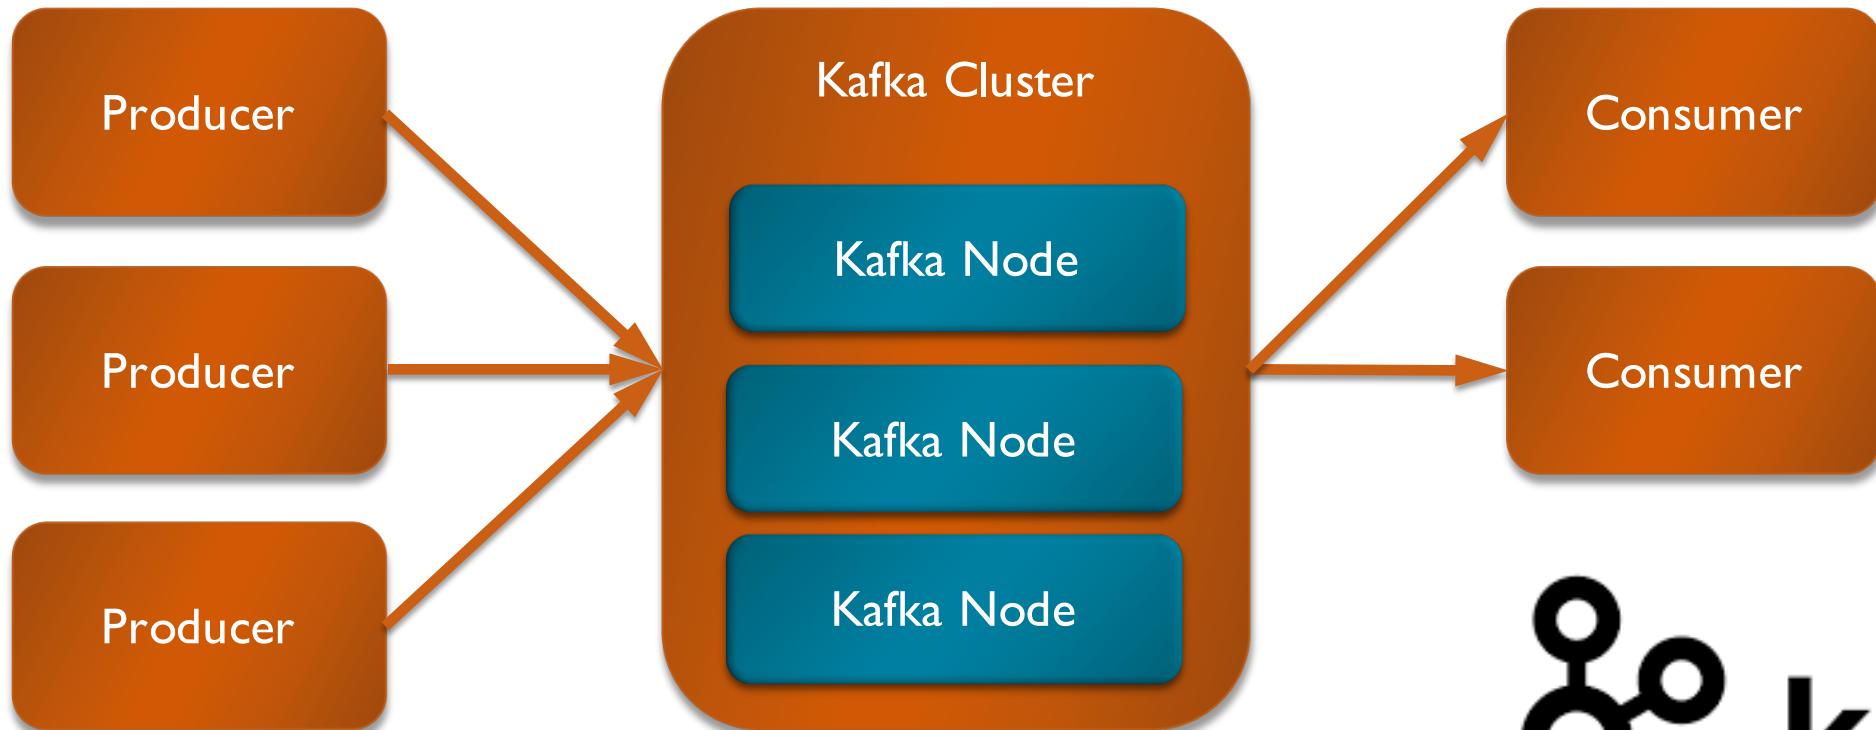


Kafka

Publish-subscribe messaging
rethought as a distributed commit log



Kafka



Kafka Use-Cases

- Messaging
- Website activity tracking
- Metrics
- Log aggregation
- Event sourcing
- **Stream processing**



Common Characteristics

Kafka, Spark & Cassandra are

- Scalable
- Distributed
- Partitioned
- Replicated



KillrPower

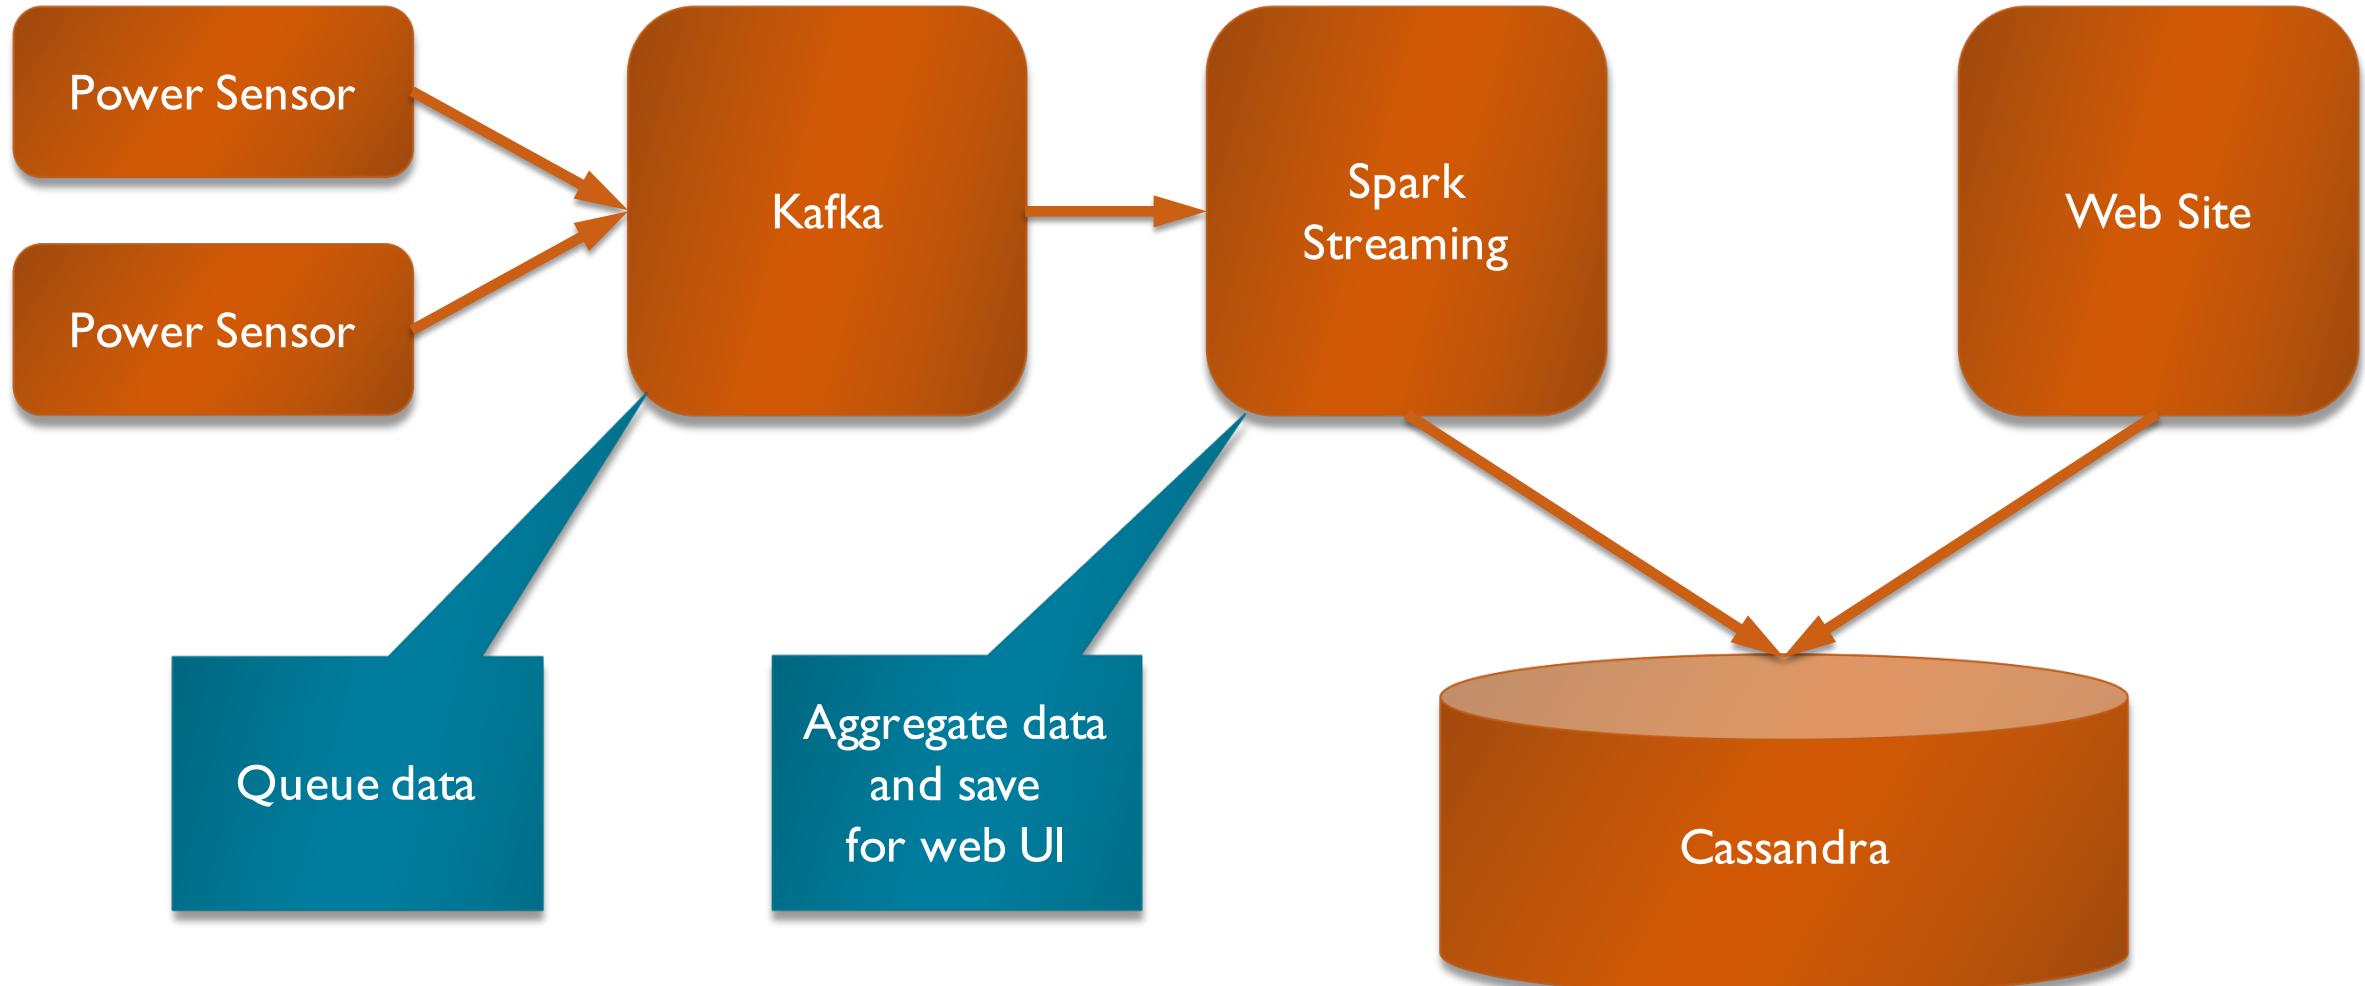
Show current power consumption in Germany
in real time

- Power consumption of all cities on a map
- Graph of data for a selected city

KillrPower – Data modeling 101

1. Collect UI Requirements
2. Model data (flow) per web UI request
3. Build Spark Streaming code to generate that data
4. Connect sensors and Spark Streaming using Kafka

KillrPower – Tech Stack



KillrPower – Tech Stack

Techs used in the demo

- Spray (HTTP for Akka Actors)
- AngularJS
- Kafka
- Spark Streaming
- Cassandra



Tips

- Push down predicates to Cassandra
- Partition your data “right”
- Think distributed



Cassandra + Spark in DSE

DSE integrates
Spark + Cassandra

[https://academy.datastax.com/
courses/getting-started-
apache-spark](https://academy.datastax.com/courses/getting-started-apache-spark)





KillrPower Live Demo



You've got the lighter, let's start a fire: Spark and Cassandra

Robert Stupp

Solutions Architect @ DataStax, Committer to Apache Cassandra

@snazy

Cassandra Days brought to you by DataStax



Thank You



BACKUP SLIDES

BACKUP - DDL



```
-- cities is just master data
CREATE TABLE cities (
    loc_id          bigint PRIMARY KEY,
    name            text,
    lat             float,
    lon             float,
    population     int,
    area            float
);
-- ... preloaded with some data
COPY cities FROM 'cities.csv';
```

```
-- Contains the timestamp of the most recent data in  
-- power_by_time and power_by_city tables below.
```

```
CREATE TABLE lastwindow (  
    dummy          text PRIMARY KEY,  
    timewindow     timestamp  
) WITH default_time_to_live = 3600;
```

BACKUP - DDL

```
-- Contains the power consumption for cities partitioned
-- by timestamp (power consumption by date/time).
CREATE TABLE power_by_time (
    timewindow          timestamp,
    --
    city_id              bigint,
    --
    lat                  float,
    lon                  float,
    name                text,
    population          int,
    area                float,
    power               float,
    --
    PRIMARY KEY ( timewindow, city_id )
) WITH default_time_to_live = 3600;
```

BACKUP - DDL

```
-- Contains the power consumption for cities partitioned
-- by city (power consumption by city).
CREATE TABLE power_by_city (
    city_id          bigint,
    --
    timewindow       timestamp,
    --
    lat              float      STATIC,
    lon              float      STATIC,
    name             text       STATIC,
    population       int        STATIC,
    area             float      STATIC,
    --
    power            float,
    --
    PRIMARY KEY ( city_id, timewindow )
) WITH CLUSTERING ORDER BY (timewindow DESC)
    AND default_time_to_live = 3600;
```

BACKUP – Spark Streaming



```
def main(args: Array[String])
{
    val conf = new SparkConf(true).
        setAppName("cstarSparkStreaming")

    val ssc = new StreamingContext(conf, slideDuration)
```

BACKUP – Spark Streaming



```
// Setup the Kafka topic Spark DStream
val records: ReceiverInputDStream[(String, Array[Byte])] = KafkaUtils.createStream[String, Array[Byte], String, String](ssc,
  kafkaParams,
  topics,
  StorageLevel.MEMORY_AND_DISK_SER_2)
```

BACKUP – Spark Streaming



```
val powerdiff = records
  // Deserialize the power record from the Kafka topic.
  .map(keyValue => deserialize[Power](keyValue._2))
  // Extract the ID of the city
  .map(power => (power.id, power))
  // Reduce by ID of the city (aggregate the power consumption)
  .reduceByKeyAndWindow((power1: Power, power2: Power) =>
    power1.copy(...),
    windowDuration,
    slideDuration)
  // Inject the current time-window timestamp
  .transform((rdd, time) => rdd.map(power =>
    {
      ...
      TimestampedPower(...)
    })
  // Map the case class into a tuple to be stored in Cassandra
  .map(...)

  // Cache the result since it's needed multiple times below
  .cache
```

BACKUP – Spark Streaming



```
// Save records per time-window
//
// That's the power consumption over the last 60 seconds (window
// for each city by time
powerdiff.saveToCassandra(cassandraKeyspace, "power_by_time",
                           SomeColumns("timewindow", "city_id",
                                       "name", "area", "populati

// Save records per city
//
// That's the power consumption over the last 60 seconds (window
// for time-window by city
powerdiff.saveToCassandra(cassandraKeyspace, "power_by_city",
                           SomeColumns("timewindow", "city_id",
                                       "name", "area", "populati
```

BACKUP – Spark Streaming



```
// Save last processed timestamp
//
// For this demo we have to save the current (means: "last
// find the current power consumption per city in the tab
powerdiff.map(tup => tup._1)
    .reduce((t, x) => t)
    // at this point we only have the one "row" containing
    // timestamp to be saved to Cassandra
    .map(t => ("dummy", t))
    .saveToCassandra(cassandraKeyspace, "lastwindow",
        SomeColumns("dummy", "timewindow"))
```

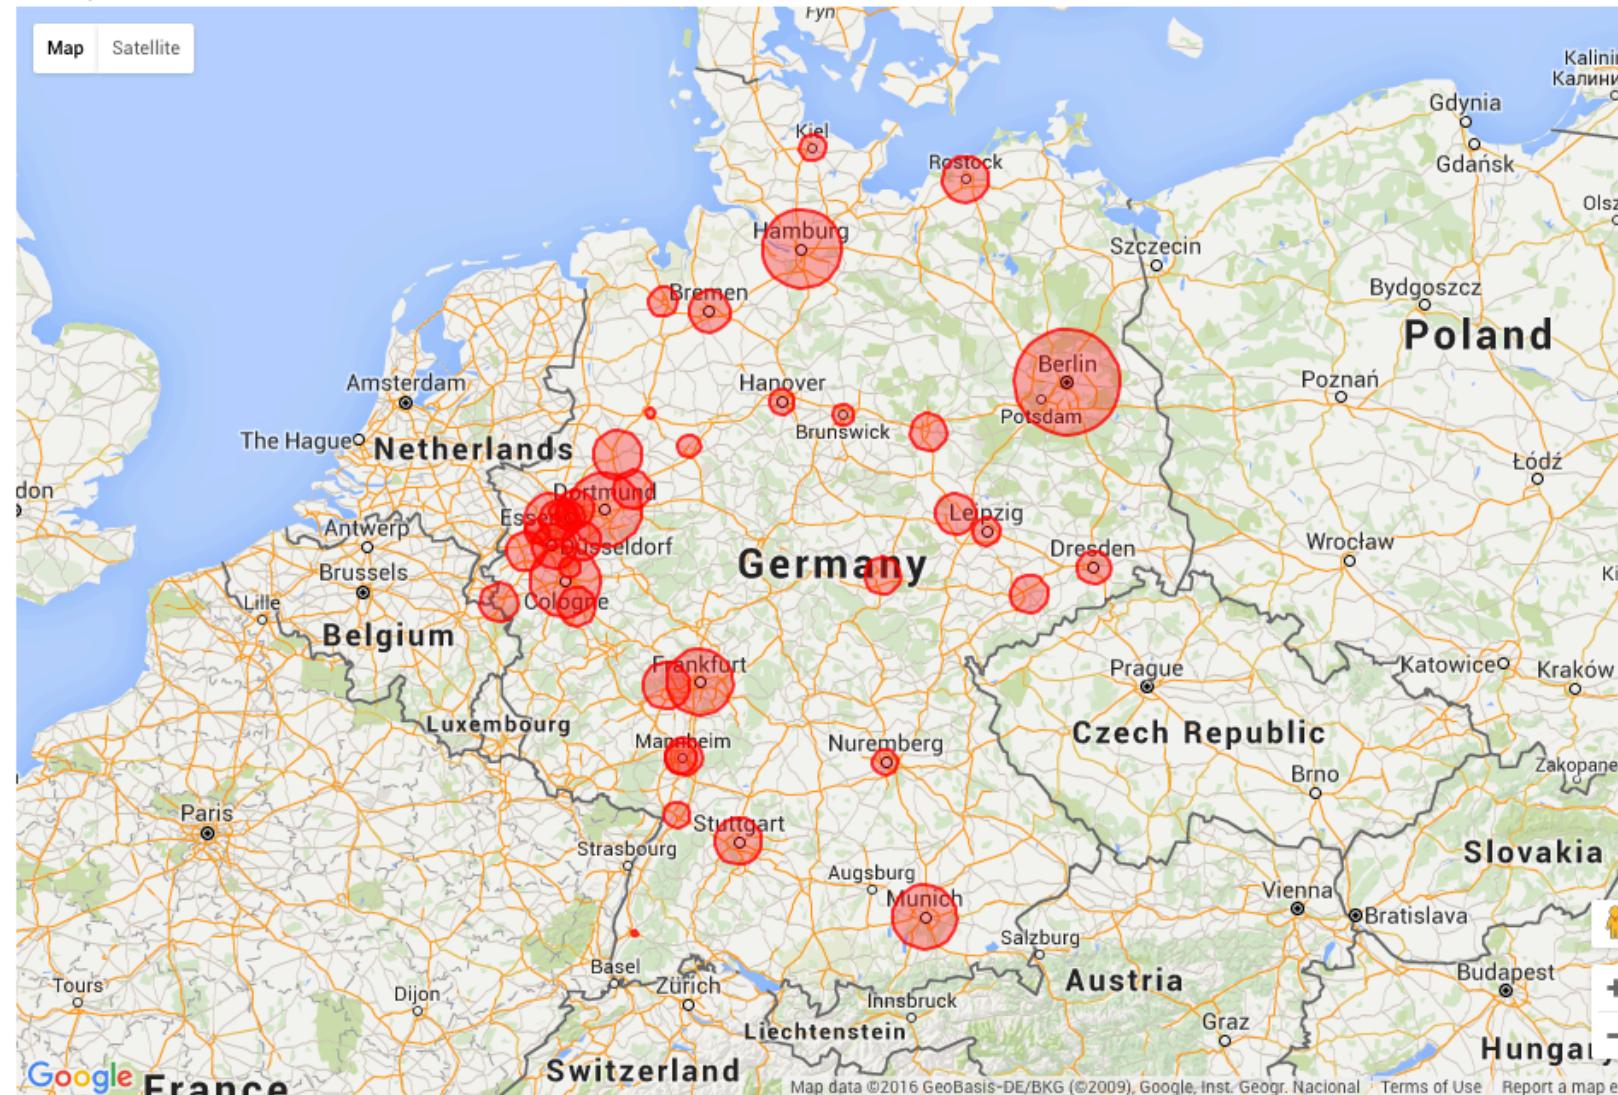
BACKUP – Spark Streaming



```
ssc.start()
```

```
ssc.awaitTermination()
```

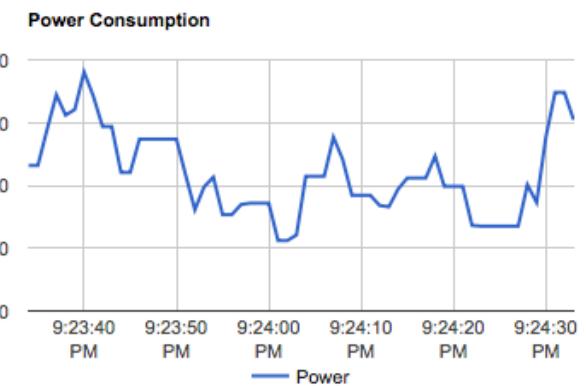
Last updated: Jan 29, 2016 9:24:33 PM



Köln

Power consumption: 456.79 kW

Population: 1,024,373
Area: 405.00 km²



This is a demo application showing fake data. Thanks to the following projects and products!

[Apache Cassandra](#) (database), [Apache Spark](#) (analytics), both combined in [DataStax Enterprise](#). Also [Apache Kafka](#) for piping data.

Coded in Scala and JavaScript using Spray, Akka, AngularJS, Google Maps & Charts